3. Deployment of Application using Kubernetes

1. Install the eksctl

```
~ $ curl -sLO https://github.com/weaveworks/eksctl/releases/latest/download/eksctl_Linux_amd64.tar.gz
~ $ tar -xzf eksctl_Linux_amd64.tar.gz
~ $ sudo mv eksctl /usr/local/bin
~ $
```

2. Check the status of eksctl

3. Create a kubernetes cluster

```
~ $ eksctl create cluster \
> --name brain-eks \
> --region ap-south-1 \
> --nodegroup-name brain-nodes \
> --node-type t3.micro \
2025-12-16 12:21:02 [ℹ]  eksctl version 0.216.0
2025-12-16 12:21:02 [ℹ]  using region ap-south-1
2025-12-16 12:21:03 [ℹ]  setting availability zones to [ap-south-1b ap-south-1c ap-south-1a]
2025-12-16 12:21:03 [ℹ]  subnets for ap-south-1b - public:192.168.0.0/19 private:192.168.96.0/19
2025-12-16 12:21:03 [ℹ]  subnets for ap-south-1c - public:192.168.32.0/19 private:192.168.128.0/19
2025-12-16 12:21:03 [ℹ]  subnets for ap-south-1a - public:192.168.64.0/19 private:192.168.160.0/19
2025-12-16 12:21:03 [ℹ]  nodegroup "brain-nodes" will use "" [AmazonLinux2023/1.32]
2025-12-16 12:21:03 [!]  Auto Mode will be enabled by default in an upcoming release of eksctl. This means managed node groups and managed net
reated by default. To maintain current behavior, explicitly set 'autoModeConfig.enabled: false' in your cluster configuration. Learn more: htt
```

4. Check cluster created or not

```
~ $ eksctl get cluster --region ap-south-1
NAME                    REGION          EKSCTL CREATED
brain-eks               ap-south-1      True
~ $
```

5. Check the nodegroup status

```
~ $ eksctl get nodegroup --cluster brain-eks --region ap-south-1
CLUSTER         NODEGROUP       STATUS  CREATED                 MIN SIZE        MAX SIZE        DESIRED CAPACITY        INSTANCE TYPE   IMAGE ID                ASG NAME        T
YPE
brain-eks       brain-nodes     ACTIVE  2025-12-16T12:32:31Z    2               2               2                       t3.micro        AL2023_x86_64_STANDARD  eks-brain-nodes-f
0cd9392-e936-5d50-cedc-5601a1366a67     managed
~ $
```

6. Check nodes status

```
~ $ kubectl get nodes
NAME                                              STATUS   ROLES    AGE     VERSION
ip-192-168-0-234.ap-south-1.compute.internal      Ready    <none>   5m57s   v1.32.9-eks-ecaa3a6
ip-192-168-87-176.ap-south-1.compute.internal     Ready    <none>   5m59s   v1.32.9-eks-ecaa3a6
~ $
```

7. Create deployment file

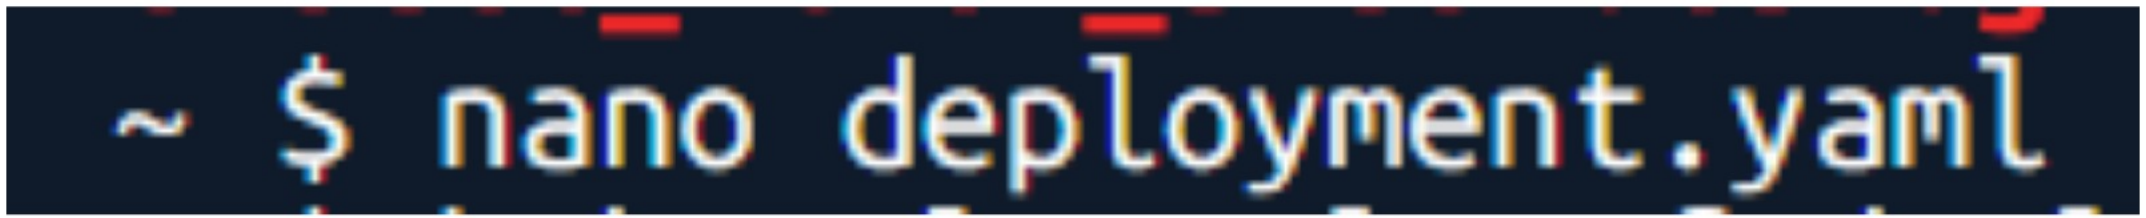

8. Write deployment code

CloudShell

ap-south-1 +

```
          "createdAt": "2025-12-15T16:19:29.386000+00:00",
  GNU nano 8.3                                                    deployment.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: brain-tasks-deployment
spec:
  replicas: 2
  selector:
    matchLabels:
      app: brain-tasks
  template:
    metadata:
      labels:
        app: brain-tasks
    spec:
      containers:
      - name: brain-tasks
        image: 284419413658.dkr.ecr.ap-south-1.amazonaws.com/brain-tasks-nginx:latest
        ports:
```

9. Scale nodes in nodegroup

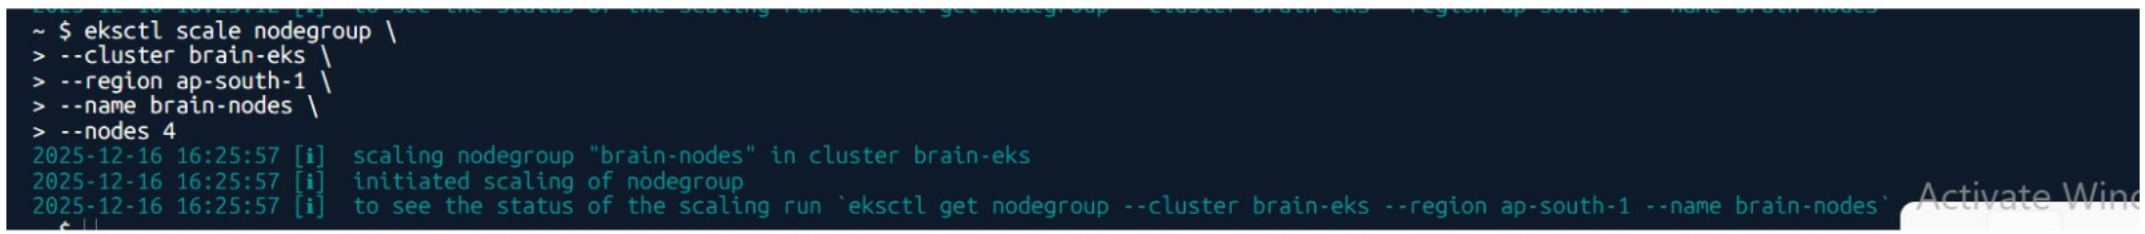

10. Check the status of nodes

```
~ $ kubectl get nodes
NAME                                              STATUS     ROLES    AGE     VERSION
ip-192-168-0-234.ap-south-1.compute.internal      Ready      <none>   3h52m   v1.32.9-eks-ecaa3a6
ip-192-168-35-76.ap-south-1.compute.internal      NotReady   <none>   4s      v1.32.9-eks-ecaa3a6
ip-192-168-59-169.ap-south-1.compute.internal     NotReady   <none>   5s      v1.32.9-eks-ecaa3a6
ip-192-168-87-176.ap-south-1.compute.internal     Ready      <none>   3h52m   v1.32.9-eks-ecaa3a6
~ $
```

11. Apply deployment

```
~ $ kubectl apply -f deployment.yaml
deployment.apps/brain-tasks-deployment created
```

12. Check deployment status

```
~ $ kubectl get deployments
NAME                      READY   UP-TO-DATE   AVAILABLE   AGE
brain-tasks-deployment    2/2     2            2           30m
~ $
```

13. Check pods running or not

```
~ $ kubectl get pods
NAME                                    READY   STATUS    RESTARTS   AGE
brain-tasks-deployment-54449f7dd-bbp7s   1/1     Running   0          31m
brain-tasks-deployment-54449f7dd-bkjbf   1/1     Running   0          31m
~ $
```

14. Create service file

```
~ $ nano service.yaml
```

15. Write service configs

ap-south-1 +

2025-12-16 16:23:12 [ℹ] to see the status of the scaling run eksctl get nodegroup --cluster brain-eks --region ap-south-1 --name brain-nodes

GNU nano 8.3 service.yaml

```
apiVersion: v1
kind: Service
metadata:
  name: brain-tasks-service
spec:
  type: LoadBalancer
  selector:
    app: brain-tasks
  ports:
  - port: 3000
    targetPort: 3000
```

16. Apply service file

```
~ $ kubectl apply -f service.yaml
service/brain-tasks-service created
~ $
```

17. Check service running or not

```
~ $ kubectl get svc
NAME                 TYPE           CLUSTER-IP      EXTERNAL-IP                                                              PORT(S)          AGE
brain-tasks-service  LoadBalancer   10.100.22.217   a192825c0868644e08bbd11a77f9276c-1755325760.ap-south-1.elb.amazonaws.com  3000:31081/TCP   31s
kubernetes           ClusterIP      10.100.0.1      <none>                                                                   443/TCP
~ $
```

18. Here website successfully runing on port 3000